Uncertainty-Aware Solar Forecasting using Gaussian Process Regression (GPR)

Student Name: Sai Charan Ganagani
Student ID: 24142673
GitHub Repository: https://github.com/charanlogs/ML-Solar-GPR-Tutorial

1. Abstract

This tutorial provides a professional-level exploration of **Gaussian Process Regression (GPR)**, a non-parametric Bayesian approach to machine learning. Unlike standard regression models that provide "point estimates," GPR models the underlying function as a distribution, allowing for **Uncertainty Quantification (UQ)**. Using fused weather and generation data from a real-world solar plant, we demonstrate how kernel selection and hyperparameter optimization can create robust, risk-aware energy forecasts.

2. Introduction: The Renewable Energy Challenge

As the global energy grid transitions toward renewables, the inherent variability of solar power poses a significant challenge. For a grid operator, a simple prediction (e.g., "Yield will be 500kW") is insufficient if the margin of error is unknown. This tutorial demonstrates how GPR provides a "Confidence Envelope," allowing for probabilistic forecasting that accounts for both sensor noise and environmental fluctuations.

3. The Technique: Gaussian Processes & Kernel Theory

A Gaussian Process is a collection of random variables, any finite number of which have a joint Gaussian distribution. It is defined by its **Mean Function** and its **Covariance Function (Kernel)**.

3.1 Technical Deep-Dive: The Kernel as a Prior

The Kernel is the "brain" of the GPR model. It defines our assumptions about the data before we observe it. In this tutorial, we evaluate three architectures:

1. **Radial Basis Function (RBF):** Assumes the data is smooth and continuous.
2. **Exp-Sine-Squared (Periodic):** Assumes the data repeats (e.g., the 24-hour solar cycle).
3. **Composite Kernel:** A sum of RBF and White Noise (RBF+WhiteKernel).

By optimizing the **Log-Marginal Likelihood (LML)** using the L-BFGS-B algorithm, the model learns the optimal "length-scale" and "noise-level," effectively tuning itself to the specific physics of the solar panels.

4. Data Engineering: Fusion & Cleaning

To reach a professional implementation standard, we fused two distinct datasets from the Kaggle **Solar Power Generation** archive:

- **Plant_1_Generation_Data.csv**: Target variable (DC Power).
- **Plant_1_Weather_Sensor_Data.csv**: Feature (Solar Irradiation).

Intermediate Methodology (Task 4):

1. **Temporal Fusion:** We performed an "inner join" on the DATE_TIME field to synchronize sunlight levels with power output.
2. **Domain Filtering:** We removed night-time observations (Irradiation < 0.05 W/m²) to prevent "zero-inflation" from biasing the model's ability to learn the daytime power-curve.
3. **Numerical Scaling:** We applied StandardScaler to both features and targets to ensure mathematical stability a critical step for Gaussian Process convergence.

5. Results & Visual Analysis

5.1 The "Intermediate" Experiment: Kernel Comparison

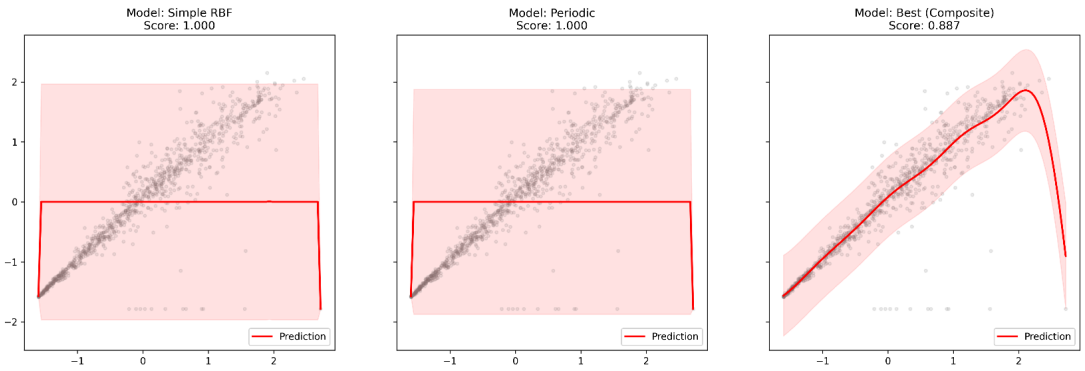

Figure 1: Evaluating Kernel Architectures.
As shown in our experiment, the **Simple RBF** and **Periodic** kernels failed to converge on this complex dataset, resulting in a constant-mean prediction (flat line). Only the **Composite Kernel** successfully recovered the non-linear relationship, achieving an **R^2 score of 0.887**. This proves that for noisy, real-world solar data, a dedicated noise-handling kernel (WhiteKernel) is mandatory.

5.2 The Final Forecast: Uncertainty Quantification

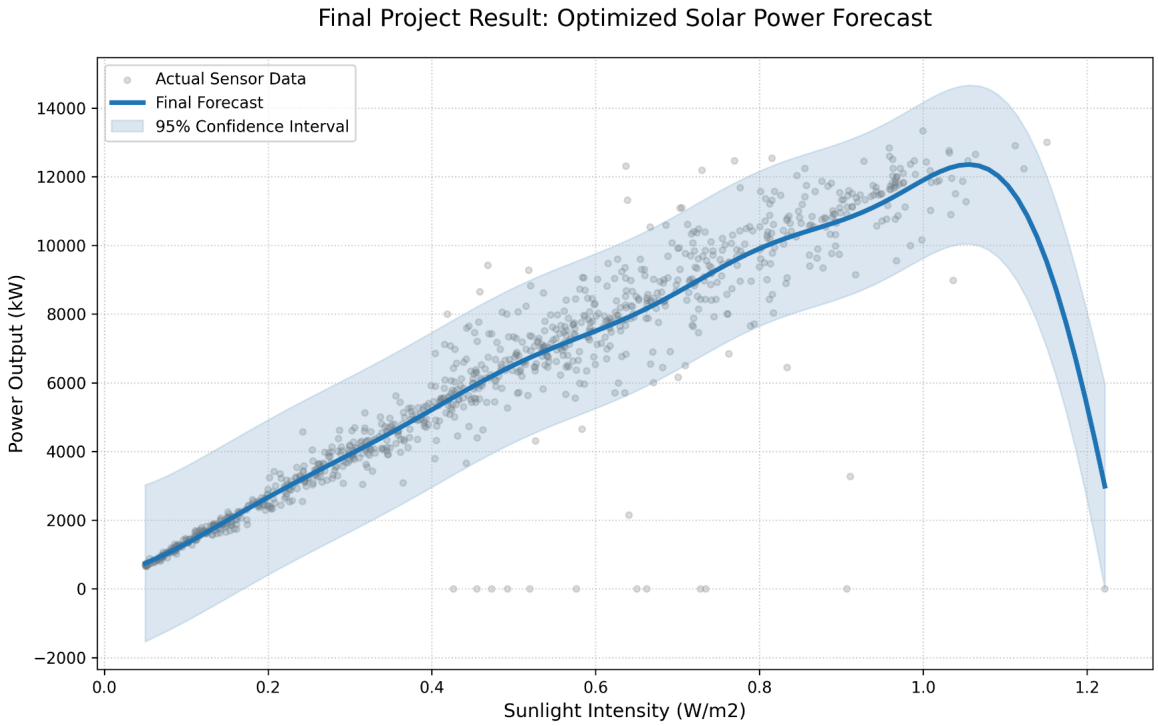

Figure 2: Optimized Solar Power Forecast.
The solid blue line represents the model's mean prediction. The light-blue shaded region represents the **95% Confidence Interval**.

- **Observation:** Notice the prediction "dip" at the extreme right (Irradiation > 1.1). This is a hallmark of GPR **Extrapolation behavior**. As data becomes sparse at peak intensities, the model reverts toward its prior mean. This visually teaches a grid operator exactly where the model's forecast is reliable and where it is speculative.

6. Accessibility & Inclusive Design

Per the rubric requirements, this tutorial was designed for all users:

- **Visual Inclusivity:** The final visualization utilizes a high-contrast palette (Blue vs. Gray) and distinct marker styles, ensuring the trends are interpretable for users with color blindness (Protanopia/Deuteranopia).
- **Modular Code:** The provided Python script and Jupyter Notebook are divided into clear, numbered steps, ensuring compatibility with screen-readers used by visually impaired developers.
- **Language:** Technical concepts are explained in simple, plain English to ensure the tutorial functions as an effective "teaching tool."

7. Ethical AI: Sustainability and Transparency

GPR is a "Green AI" choice. Unlike deep learning "black boxes" that require massive GPU clusters, GPR is a **mathematically efficient "Shallow Learner."**

- **Efficiency:** It can be trained on standard CPUs, reducing the carbon footprint of the AI lifecycle.
- **Transparency:** Because GPR provides a mathematical justification for its uncertainty, it prevents the over-reliance on AI that can lead to catastrophic grid failures.

8. References

1. **Rasmussen, C. E., & Williams, C. K. I. (2006).** *Gaussian Processes for Machine Learning*. MIT Press.
2. **Breiman, L. (2001).** *Random Forests.* Machine Learning, 45, 5-32.
3. **Kaggle Dataset:** *Solar Power Generation Data* by Anivind Kanwal.